

INNOPOLIS UNIVERSITY
Big Data — Final Project, 2026

Detecting IoT Malware at Network Scale

Team: Team 30
Repository: [Dmitry057/big-data-group-project](#)
Local path: /home/team30/project30

Ivan Ilyichev	Data Engineer
Kirill Shumsky	ML Specialist
Nikita Zagainov	Data Scientist
Dmitry Tetkin	Tester & Technical Writer

May 12, 2026

1 Introduction

1.1 Why we picked this problem

Internet-of-Things devices — cheap routers, IP cameras, DVRs — ship with default passwords and almost never get patched. Once compromised, they are pulled into botnets like *Mirai*, *Hide-and-Seek* and *Torii*, and used to generate DDoS attacks, port scans and command-and-control traffic. The malicious flows look almost identical to the device’s own normal traffic, and there are millions of them per hour on a real network. So we wanted a system that can sit in front of that firehose, look at every connection, and say: *this one is fine, this one is part of a port scan, this one is a DDoS amplifier*.

The IoT-23 dataset gave us labelled, real-world data to do exactly that — roughly 25 million labelled connections from 12 IoT-device captures. Concretely, the goals of the project are:

1. Move 25M Zeek connection records from CSV into a Hive warehouse on HDFS using the standard Hadoop tools.
2. Run a small set of HiveQL queries to understand what’s actually in the data — how the classes split, which ports get attacked, what the typical packet shapes look like.
3. Train a distributed classifier in Spark that tells the seven classes apart, evaluate it on a held-out set, and put everything in a Superset dashboard the rest of the team can read.

The whole thing is wired into a single `main.sh` so the grader can re-run every stage from scratch.

1.2 What’s in this document

- **Data Description** — where the data comes from, how big it is, and how the four pipeline stages connect.
- **Data Preparation** — the relational schema, how we loaded it into Postgres and lifted it into Hive, and the storage-format benchmark.
- **Data Analysis** — seven HiveQL insights, with the chart and a short interpretation for each.
- **ML Modeling** — the Spark feature pipeline, the two models we trained, and how they did on the test set.
- **Data Presentation** — what’s on the Superset dashboard and what it tells the viewer.
- **Conclusion** — what we shipped, what gave us trouble, what we’d do differently, and the contribution table.

1.3 Pipeline at a glance

Stage	Script	Output
1 — Ingestion	<code>stage1.sh</code>	PostgreSQL DB + HDFS Parquet warehouse
2 — Storage & EDA	<code>stage2.sh</code>	Hive partitioned tables + 7 EDA result tables
3 — Predictive analytics	<code>stage3.sh</code>	Trained Spark ML models + held-out predictions
4 — Dashboard	<code>stage4.sh</code>	Apache Superset dashboard refreshed

Table 1: The four stages of the pipeline.

2 Data Description

2.1 Where the data comes from

We used a hand-picked subset of **Aposemat IoT-23** ([Kaggle Source](#), the full dataset is available [1]), a public dataset released by the Stratosphere Lab at Czech Technical University. It's a collection of network captures recorded on real IoT devices that were either left in their normal state or deliberately infected with a known malware family. Each capture comes as a Zeek `conn.log`, which is basically one row per network connection with both a binary label (**Benign/Malicious**) and a more specific label for the malicious ones (DDoS, C&C, port-scan, etc.).

We took 12 of these captures — enough to cover several malware families without making the warehouse unmanageable on a 3-node student cluster.

2.2 How big is it

Property	Value
Total connections (rows)	25,011,003
Distinct captures	12
Per-row features (raw Zeek fields)	23
Target classes	7
Format on disk (cluster)	Parquet + Snappy
Approx. raw CSV size	~3.39 GB

Table 2: Volume and shape of the working dataset.

The seven label values are: **Benign**, **Malicious**, **Malicious DDoS**, **Malicious PartOfAHorizontalPortScan**, **Malicious C&C**, **Malicious Attack**, **Malicious FileDownload**.

2.3 Architecture of the pipeline

Figure · Pipeline architecture

From a folder of CSVs to a refreshing dashboard.




Figure 1: End-to-end architecture. CSVs go into Postgres; Sqoop lifts the typed tables onto HDFS as Parquet; Hive serves both the EDA queries and Spark; the trained model and the EDA result tables both feed a single Superset dashboard.

2.4 Per-stage input and output

Stage	Input	Output
Stage 1	12 IoT-23 capture CSV files	PostgreSQL tables captures, network_connections; HDFS Parquet+Snappy warehouse at /user/team30/project/warehouse
Stage 2	Sqoop Parquet warehouse on HDFS	Hive database team30_projectdb with bucketed and partitioned tables; 7 result tables q1_results...q7_results; CSV exports in output/
Stage 3	Hive table network_connections_part (label-partitioned)	Trained models models/model1 (RF), models/model2 (LR); predictions output/modelN_predictions.csv; metrics output/evaluation.csv
Stage 4	Hive result tables + evaluation.csv	Apache Superset dashboard

Table 3: Inputs and outputs of each pipeline stage.

3 Data Preparation

3.1 The relational model

In Postgres we have two tables and one foreign key. `captures` holds metadata about each of the 12 IoT-23 captures (one row per capture). `network_connections` holds every parsed Zeek connection (~25M rows) and points back to its capture via `capture_id`. That's the whole schema — nothing fancy.

Figure · Relational schema

Two tables, one foreign key — all the structure we need.




Figure 2: ER diagram. `captures` (1) — (N) `network_connections`. The crow's-foot on the right means *many connections per capture*.

To load the CSVs we go through a third table, `stg_network_connections`, which is just a staging area where every column is plain text. We `COPY` the CSV in as-is, then a transformation `INSERT` cleans the cells (regex out the trailing `.0` on integers, swap sentinel `-` for `NULL`, cast to `inet` and `timestampz`) and writes the result into the typed `network_connections` table. Doing it in two steps means we never have to fight the Postgres CSV parser over IoT-23's quirks.

3.2 Sample records

A couple of rows from `network_connections` (truncated for width):

cap	ts	id_orig_h	proto	conn_state	orig_bytes	resp_bytes	orig_pkts	label
1	2018-12-21 14:02:11	192.168.1.132	tcp	S0	0	0	1	Malicious
9	2019-04-30 09:15:48	192.168.100.103	udp	SF	1024	64	4	Benign

Table 4: Two example records from the canonical `network_connections` table.

3.3 The Hive warehouse

The Hive database is called `team30_projectdb` and lives on top of the Sqoop output. We keep two physical layouts:

- `captures_bucketed` — 12 rows, but `CLUSTERED BY (capture_id) INTO 4 BUCKETS` so it joins cleanly against the fact table.
- `network_connections_part` — partitioned by `label`. Almost every query we ever run filters by `label`, so partitioning that way means Tez (and later Spark) can skip whole files instead of scanning everything.

The whole warehouse is rebuilt by `sql/db.hql`:

```
SET hive.execution.engine=tez;
SET hive.exec.dynamic.partition=true;
SET hive.exec.dynamic.partition.mode=nonstrict;
SET hive.enforce.bucketing=true;

INSERT OVERWRITE TABLE network_connections_part PARTITION (label)
SELECT connection_id, capture_id, CAST(ts AS TIMESTAMP), uid,
       id_orig_h, id_orig_p, id_resp_h, id_resp_p,
       proto, service, duration, orig_bytes, resp_bytes,
       conn_state, local_orig, local_resp, missed_bytes, history,
       orig_pkts, orig_ip_bytes, resp_pkts, resp_ip_bytes,
       tunnel_parents, detailed_label, label
FROM network_connections_ext;
```

3.4 Picking a storage format

Before settling on Parquet + Snappy we ran a small benchmark: write the fact table four different ways, then measure the on-disk footprint (via `hadoop fs -du`) and time a full scan and one analytic query. Three runs each, results averaged:

Format	Codec	Runs	Size (B)	Avg write (s)	Avg full scan (s)	Avg analytic (s)
parquet	snappy	3	4,506,447	17.7	6.76	9.29
parquet	gzip	3	3,003,399	18.0	6.64	9.56
avro	snappy	3	4,941,651	17.0	9.32	10.50
avro	bzip2	3	3,082,133	19.0	8.36	11.01

Table 5: Storage format benchmark. Parquet beats Avro on scan time by a clear margin. On size, gzip and bzip2 shave roughly 33% off Snappy — a real saving on disk but at a CPU cost on every read; Snappy keeps scans the fastest at a sensible footprint, so we shipped Parquet + Snappy. Raw numbers in `output/stage1_storage_benchmark_avg.csv`.

4 Data Analysis

The seven HiveQL queries (`sql/q1.hql...sql/q7.hql`) each materialise a small `qN_results` Parquet table and dump a CSV that the dashboard reads directly. The interpretations below are what we'd actually tell a teammate looking at the chart for the first time.

4.1 Insight 1 — How the labels split




Figure 3: Class shares of `network_connections_part`.

Label	Rows	Share
Benign	8,780,158	35.11%
Malicious	7,055,007	28.21%
Malicious DDoS	5,778,154	23.10%
Malicious PartOfAHorizontalPortScan	3,386,241	13.54%
Malicious C&C	8,685	0.035%
Malicious Attack	2,755	0.011%
Malicious FileDownload	3	0.000%

Table 6: Label distribution.

The dataset isn’t balanced — four classes hold 99.9% of the rows — but it’s not flat either. The three rare classes (C&C, Attack, FileDownload) are interesting precisely *because* they’re rare: collapsing them into a single “malicious” bucket would lose a real signal. That’s why we kept the task multiclass instead of going binary.

4.2 Insight 2 — Which protocol gets used




Figure 4: Protocol usage per label. UDP shows up almost exclusively in benign DNS-style traffic; almost every malicious connection is TCP.

If you see a UDP flow on this network, it's almost certainly fine — DNS, NTP, the usual. If you see a malicious flow, it's basically always TCP. So `proto` on its own won't separate one malicious class from another, but it's a strong prior against “UDP = bad”.

4.3 Insight 3 — How big the connections are




Figure 5: Average originator and responder bytes per connection by label.

Label	Avg orig bytes	Avg resp bytes	Avg total bytes
Benign	12.87	3,830.63	3,843.51
Malicious	0.54	1.87	2.41
Malicious DDoS	3.71×10^8	0.60	3.71×10^8
Malicious C&C	31.48	517.93	549.41
Malicious Attack	4.35	6.16	10.51
Malicious PartOfAHorizontalPortScan	0.87	0.00	0.87
Malicious FileDownload	83.33	64,982.00	65,065.33

Table 7: Average per-connection byte counts by label.

DDoS jumps out immediately — 3.7×10^8 originator bytes per flow on average, basically six orders of magnitude over benign. That’s the amplification payload. At the other end, port scans and the generic **Malicious** class barely move a byte per connection: they’re basically just SYN probes. C&C looks more like a normal long-lived session (a few KB on the responder side), and FileDownload (only 3 rows) shows the asymmetric pattern you’d expect.

4.4 Insight 4 — Each capture is one behaviour




Figure 6: Total flows and malicious ratio per capture.

The captures don’t sit anywhere in the middle: 8 of them are over 79% malicious, 4 of them are under 11%. There’s no “slightly suspicious” capture. That makes sense — each capture was recorded around a single infection, so basically the whole capture is either compromised traffic or a clean baseline.

The practical consequence: if we ever do a real evaluation, we should split by capture, not by row. Random row-level splitting (which is what we used for this report) leaks the same capture into both train and test. We flagged it in the conclusion.

4.5 Insight 5 — What the connection state looks like

Almost everything in this dataset is **S0** — a SYN that never gets answered. Even the benign traffic is dominated by IoT devices repeatedly trying to reach a server that’s not there. DDoS is the exception: it splits between **OTH** (we joined the flow mid-stream) and **RSTOS0** (the originator gets a RST after the SYN). C&C shows the classic **S3** fingerprint — the originator finishes but the responder is still talking, which is what a long-running control channel looks like.

Label	Dominant conn_state	Rows	Share within label
Benign	S0	8,722,217	99.3%
Malicious	S0	7,036,199	99.7%
Malicious DDoS	OTH + RSTOS0	5,777,712	99.99%
Malicious C&C	S0 + S3	7,655	88%
Malicious PartOfAHorizontalPortScan	S0	3,386,241	100%

Table 8: Dominant TCP connection state per label.

4.6 Insight 6 — Where the attacks are aimed




Figure 7: Top targeted destination ports per non-Benign label.

Telnet (port 23) is, by a huge margin, the most attacked destination — 4M generic Malicious flows and another 3.4M port-scan flows hit it. That’s the Mirai pattern: scan the internet for IoT devices with default Telnet credentials, log in, install yourself, repeat. DDoS prefers the high port 62336 (a Mirai control/echo port) and HTTP 80. C&C mostly uses IRC port 6667, which is also a Mirai-family classic.

Label	Port	Flows
Malicious	23	4,055,327
Malicious DDoS	62336	3,578,391
Malicious PartOfAHorizontalPortScan	23	3,386,119
Malicious	81	2,571,979
Malicious DDoS	80	1,125,806
Malicious DDoS	992	1,008,351
Malicious C&C	6667	6,706

Table 9: Top 7 targeted destination ports across non-Benign labels.

4.7 Insight 7 — How many packets per connection




Figure 8: Average originator and responder packets per connection by label.

The interesting row here is C&C: about 495 packets out per session, with a real (8-packet) response. That's a genuine long-running conversation, very different from the one-shot SYN of port scans (where the responder is essentially zero). The combination of `orig_pkts`, `resp_pkts` and `conn_state` ends up doing most of the work for the Random Forest later.

5 ML Modeling

5.1 Building the feature pipeline

`scripts/ml_data_preparation.py` fits a single PySpark Pipeline on the training fold (so the encoders never see the test rows). Five steps:

1. **Cyclical time encoding.** A small custom transformer pulls year, month, day, hour, minute, second out of `ts` and adds sin/cos for the periodic ones. Without this the model would learn that `hour=23` and `hour=0` are far apart.
2. **Imputation.** Drop `tunnel_parents` (always null); fill string nulls with "unknown", numeric nulls with 0, booleans with `False`.
3. **Categorical encoding.** `StringIndexer` + `OneHotEncoder` for `proto`, `service`, `conn_state`, `history`.
4. **Feature selection.** A `ChiSqSelector(fpr=0.1)` on the categorical block keeps only the dummies that actually correlate with the label.
5. **Assemble + scale.** `VectorAssembler` concatenates everything; `StandardScaler(withMean=False, withStd=True)` gives us the final `features` column.

The label gets its own `StringIndexer`. We split 60/40 with seed 42, and write the resulting train/test out to HDFS as Parquet so training and evaluation read from the same on-disk artefact.

5.2 Key optimisations that made it fit on the cluster

Three things turned a slow, fragile training loop into something we could re-run reliably inside our YARN budget. They are not glamorous, but they are the reason the ML stage finished at all on a contended cluster.

- **Persisting the training set in memory + disk.** A `CrossValidator` re-reads the training fold once per hyperparameter combination and once per fold. Without caching, every one of those reads goes back to HDFS, decodes Parquet, re-runs the feature pipeline, and shuffles — which is most of the wall-clock time. We call `train_data.persist(StorageLevel.MEMORY_AND_DISK)` after the initial read so the second and subsequent passes hit RAM (or local disk if a partition spills). This alone cut the per-fit time roughly in half on our profile and let the CV runs survive single-executor losses without falling back to HDFS.
- **Selecting categorical features with `ChiSqSelector`.** One-hot encoding `proto`, `service`, `conn_state` and `history` produces several hundred sparse dummy columns, most of which carry no signal for the label. Feeding all of them into Random Forest blows up split-search time, and into Logistic Regression blows up the optimiser's inner loop. We run `ChiSqSelector(fpr=0.1)` over the categorical block, fitted on the training fold only, and keep only the dummies whose chi-squared p-value clears the threshold. The feature vector shrinks by an order of magnitude and the model fits faster without measurable loss on the held-out F1.
- **Writing train/test as Parquet with explicit sharding.** Once the feature pipeline is fitted, we materialise the transformed train/test sets back to HDFS:

```
train_data.select("features", "label") \
    .coalesce(4) \
    .write.mode("overwrite") \
    .format("parquet") \
    .save("project/data/train")
```

Two things matter here. First, we write Parquet so the training script can re-read the already-vectorised features without re-running the feature pipeline (it would otherwise do that on every `spark-submit`). Second, the explicit `coalesce(4)` controls the number of output files: too many and the executor open-file overhead dominates; too few and we lose parallelism on the read. Four shards matched our executor layout (6 executors $\times$ 3 cores) reasonably well for both training and evaluation. Combined with persistence above, the next `ml_models_training.py` run starts from a vectorised, sharded Parquet file and goes straight into the cross-validator.

5.3 Two models, one cross-validator

We trained two models, both wrapped in `CrossValidator(numFolds=2, parallelism=3)` optimising multiclass `f1`. Each `spark-submit` ran on YARN with 6 executors $\times$ 3 cores $\times$ 6 GB. The training set is `persist(MEMORY_AND_DISK)`'d so the CV folds don't repeat the read.

Model	Grid	Best
Random Forest	numTrees $\in$ {30, 60} maxDepth $\in$ {4, 6} impurity = gini (default)	numTrees = 60 maxDepth = 6 impurity = gini
Logistic Regression (multinomial)	regParam $\in$ {0.01, 0.1} elasticNetParam $\in$ {0.0, 0.5} maxIter = 50, family = multinomial	regParam = 0.01 elasticNet = 0.0

Table 10: Hyperparameter grid and the values picked by 2-fold cross-validation on the multiclass `f1` metric.

5.4 How they did

The evaluation script (`scripts/ml_evaluation.py`) reads the HDFS-persisted predictions and computes four metrics with PySpark's `MulticlassClassificationEvaluator`:

Model	Accuracy	Precision	Recall	F1
Random Forest	0.9974	0.9974	0.9974	0.9973
Logistic Regression	0.9943	0.9939	0.9943	0.9941

Table 11: Held-out evaluation on $\sim$ 10M test rows.

Both models land in the 0.99+ range across every metric. The Random Forest is slightly ahead — about 0.003 F1 over Logistic Regression — which is what we'd expect on a feature space where `conn_state` + `id_resp_p` + `orig_pkts` contain sharp non-linear thresholds that a tree ensemble can split on directly. The multinomial LR closes most of the gap once the `ChiSqSelector` block trims uninformative dummies and the `StandardScaler` brings the per-feature scales together, so the linear baseline is genuinely competitive on this problem rather than a straw man.

We persist the predictions to HDFS as CSV before evaluating, which means we can re-run the evaluator without re-running the trainer — useful if we want to add a new metric later.

6 Data Presentation

6.1 The dashboard

The dashboard has six sections:

- A. Data characteristics.** Headline numbers: 25M rows, 23 features, 7 classes, 12 captures. The first thing a viewer sees.
- B. Class balance (Insight 1).** Donut and percentage list, sourced from `q1_results`.
- C. Behaviour signatures (Insights 2, 3, 5, 7).** Stacked-bar protocol mix, average bytes/-packets per label, dominant `conn_state` per label.
- D. Top targets (Insight 6).** Port histograms per non-Benign label, with a callout on port 23.
- E. Capture strata (Insight 4).** Per-capture maliciousness ratio bar chart — the bimodal capture-level pattern is hard to miss.
- F. Model performance.** Side-by-side comparison of the two models, sourced from `evaluation.csv`.

6.2 What the dashboard tells you

Three things stand out once everything is in front of you:

1. **The class skew is interesting, not degenerate.** Four classes carry the bulk of the data and three rare classes carry real diagnostic signal. Treating the task as multiclass keeps that signal visible.
2. **The attack surface is concentrated.** Port 23 alone accounts for more than 7M malicious flows. Rate-limiting inbound Telnet on the IoT subnet would have killed most of the malicious volume in this corpus on its own.
3. **Both models land above 0.99 F1, RF slightly ahead.** With four hyperparameter combinations and a 2-fold CV, the Random Forest reaches $F_1 = 0.997$ and the multinomial Logistic Regression reaches $F_1 = 0.994$ on the same feature pipeline. The roughly 0.003 gap is consistent with what you'd expect when most of the signal lives in sharp non-linear thresholds (`conn_state`, `id_resp_p`, `orig_pkts`) — a tree ensemble splits on them directly, while LR has to approximate them through one-hot dummies. Either model is a defensible starting point for a production version.

7 Continuous Integration

The repo is wired up with a three-layer test strategy. We deliberately keep heavyweight cluster work out of the per-commit loop — the full pipeline downloads ~3.39 GB of CSVs through `gdown` and takes 30+ minutes on a contended YARN queue — and push it into a manual stage instead.

Layer	Where	Trigger	Covers
1 — Lint / unit	GitHub ubuntu-latest	Actions, every push and pull_request	ShellCheck on every *.sh; ruff + advisory pylint on scripts/*.py; pdflatex build of report/report.tex; the tests/unit/*.bats suite (repo- structure + artefact-schema invariants).
2 — Cluster smoke	hadoop-01 (the team's edge node)	manual: bash scripts/run_tests.sh smoke	Postgres row counts + FK integrity + label coverage; HDFS warehouse sizes; Hive table inventory and par- titions; q1.hql re-run via beeline with a row-count diff against output/q1.csv; Spark ml_evaluation.py re-run over persisted predictions (no retrain- ing); optional public-dashboard ping.
3 — Full pipeline	hadoop-01	manual: bash main.sh	The end-to-end run from stage1.sh through stage4.sh, including the gdown download. Intentionally not in CI.

Table 12: Three test layers, each with its own trigger and cost profile.

7.1 What the cloud workflow checks

.github/workflows/lint.yml fans out into four parallel jobs on ubuntu-latest:

- **ShellCheck** on every tracked *.sh outside the vendored tests/bats/ tree, with a small set of project-wide excludes (SC1091, SC2155, SC2034).
- **ruff + pylint** on scripts/. ruff is enforcing with rule set E,F,W; pylint runs in advisory mode (the score is reported but does not gate the build).
- **pdflatex build** — converts the SVG figures to PDF with rsvg-convert, runs pdflatex + bibtex + pdflatex twice, then uploads the built report.pdf as a workflow artefact. Catches broken references and missing figures before they reach a reader.
- **bats unit tests** — runs bash scripts/run_tests.sh unit, which executes the vendored bats-core over tests/unit/*.bats. These tests assert that the right files exist on disk and that the committed output/evaluation.csv still has Random Forest at $F1 \geq 0.99$ across all seven classes.




Figure 9: All four jobs green on the current `main` commit (`Build report.pdf` 1m, `ShellCheck` 17s, `bats unit tests` 6s, `ruff + pylint` 33s). This is the contract we get on every `push` and `pull_request`.




Figure 10: A failing intermediate run on the same PR — `ruff + pylint` caught a regression and the workflow surfaced it before the broken commit could be merged. The other three jobs still ran in parallel, so we got partial feedback in seconds.

7.2 The cluster runner

`scripts/run_tests.sh` is the single entry point for both unit-style tests (the same ones CI runs) and the cluster smoke suite:

```
# Cloud-safe checks only (what CI runs):
bash scripts/run_tests.sh unit

# Cluster smoke suite (Postgres + Hive + Spark + HDFS):
bash scripts/run_tests.sh smoke

# A single suite by filename prefix:
bash scripts/run_tests.sh smoke 02_postgres
```

Every cluster test is gated by `require_cmd` and `require_secret` helpers, so the same suite can run anywhere — on the cluster it actually validates the warehouse, and on a developer laptop the cluster tests simply skip with a clear reason.

```

All selected tests passed.
team30@hadoop-01:~/project30$ bash scripts/run_tests.sh

=== Running unit suite (2 file(s)) ===
1..15
ok 1 main.sh exists at repo root
ok 2 all four stage scripts exist and start with bash shebang
ok 3 every stage script uses set -euo pipefail
ok 4 all 7 HQL EDA queries are present
ok 5 Hive warehouse DDL exists
ok 6 PostgreSQL DDL and test queries exist
ok 7 ML pipeline scripts are all present
ok 8 .gitignore excludes data/, secrets/, .venv/
ok 9 secrets/ contents are not tracked by git
ok 10 evaluation.csv has both models and four metrics
ok 11 Random Forest F1 in evaluation.csv is above 0.99 (regression guard)
ok 12 all 7 EDA CSVs exist and are non-empty
ok 13 all 7 EDA charts exist and are non-empty
ok 14 storage-benchmark CSV is committed
ok 15 model directories are committed for both models

=== Running smoke suite (7 file(s)) ===
1..20
ok 1 output/ has all required Stage-3 artifacts
ok 2 model1_predictions.csv contains all seven labels
ok 3 presentation PDF is non-empty
ok 4 captures table has 12 rows # skip required command not on PATH: psql
ok 5 network_connections has ~25M rows (regression guard, > 24M) # skip required command not on PATH: psql
ok 6 all 7 label values present in network_connections # skip required command not on PATH: psql
ok 7 every connection references an existing capture (FK integrity) # skip required command not on PATH: psql
ok 8 no leftover '-' sentinel in id_orig_h after typed insert # skip required command not on PATH: psql
ok 9 sql/test_database.sql runs cleanly end-to-end # skip required command not on PATH: psql
ok 10 HDFS warehouse directory exists
ok 11 captures table directory exists on HDFS
ok 12 network_connections table is non-trivial on HDFS (> 500 MB)
ok 13 train/test parquet are present in project/data
ok 14 Hive database team30_projectdb is reachable
ok 15 all expected Hive tables exist
ok 16 network_connections_part is partitioned by label
ok 17 q1_results table is populated (> 0 rows)
ok 18 q1_results row count matches committed output/q1.csv
ok 19 ml_evaluation.py re-runs successfully on YARN
ok 20 public dashboard returns HTTP 200 anonymously # skip DASHBOARD_URL not set

All selected tests passed.
team30@hadoop-01:~/project30$

```

Figure 11: Layer 2 in action — `bash scripts/run_tests.sh` on the cluster from `team30@hadoop-01`. The 15 cloud-safe unit tests pass, and 13 of the 20 cluster smoke tests run green; the 7 `skip` lines are correct `require_cmd/require_secret` behaviour — `psql` is not installed on the edge node and `DASHBOARD_URL` is unset for this run. The smoke suite covers HDFS warehouse, Hive table inventory + partition layout, the `q1_results` row-count check against the committed `output/q1.csv`, and the Spark `ml_evaluation.py` re-run on persisted predictions.

7.3 What it caught

A few concrete examples from the project where CI did real work:

- Two `ml_models_training.py` commits introduced `ruff` F401 unused-import warnings that nobody noticed locally — caught by the `python-lint` job on the next push.
- A broken `\cite{}` key in `report.tex` slipped through local edits; the `latex-build` job failed loudly on the next push instead of leaving us with a quietly mis-rendered PDF.
- Several stage scripts were missing `set -euo pipefail` at one point; the `tests/unit/01_repo_structure.b` invariant catches that on `ubuntu-latest` in under five seconds.

7.4 What we shipped

A reproducible Hadoop pipeline that takes 25M Zeek connection records, loads them into a Hive warehouse on HDFS, runs seven HiveQL queries on Tez to surface the structure of the data, trains two Spark ML models with cross-validation on YARN, and exposes everything in a public Apache Superset dashboard. Both models clear $F_1 = 0.994$ on the held-out test set of about 10M rows, with Random Forest at $F_1 = 0.997$. Every artefact in `output/` can be regenerated by running `main.sh`.

7.5 How we worked

The role split worked well. Ivan owned ingestion and the Hive layout; Kirill owned the Spark feature pipeline and the model grid; Nikita owned EDA and the dashboard story; Dmitry owned documentation and ran the re-run checks. The “one `main.sh`, every stage idempotent” rule was probably the single most useful constraint — it forced us to clean up before doing anything, instead of relying on whatever was left from the last run.

7.6 What gave us trouble

- **Loading the CSVs.** The raw Zeek files use `-` for missing values and write integers as floats (1234.0). The Postgres transformation insert had to handle both with `NULLIF + REGEXP_REPLACE` before casting. One missed regex meant a full reload.
- **Partition skew in Hive.** Partitioning by `label` is great for queries but makes one tiny partition (FileDownload, 3 rows) and four huge ones. Fine for analytical scans, painful for Spark shuffles — we worked around it with `repartition(18)` after the read.
- **Tightening the LR baseline.** An early Logistic Regression pass barely beat majority-class until we routed it through the same `ChiSqSelector(fpr=0.1)` block and `StandardScaler(withStd=true)` as Random Forest. The fix wasn’t more tuning of LR, it was making sure the feature pipeline fed it inputs that were comparable in scale and free of the high-cardinality one-hot noise — which lifted F1 from ≈ 0.18 to 0.994.

7.7 What we’d do next

1. Split by `capture_id` instead of by row. Insight 4 basically tells us we have to.
2. Add Gradient-Boosted Trees to the grid. With both RF and LR already at 0.99+, GBT is the natural next ensemble to try for the last percent.
3. Wire the trained Random Forest into Spark Structured Streaming over the same Hive partitions, so it can score new traffic in near real time.—

7.8 Who did what

Task	Ivan I.	Kirill S.	Nikita Z.	Dmitry T.	Avg hrs	Deliverables
Dataset selection & access	25%	25%	25%	25%	1	IoT-23 captures (Kaggle Source)
Stage 1 — PostgreSQL schema & load	100%	0%	0%	0%	5	create_tables.sql, build_projectdb.py
Stage 1 — Sqoop & HDFS warehouse	100%	0%	0%	0%	8	import_to_hdfs.sh
Stage 1 — Storage benchmark	80%	0%	10%	10%	4	benchmark_storage_*, CSVs
Stage 2 — Hive warehouse setup	90%	0%	0%	10%	5	db.hql, partitioned tables
Stage 2 — 7 EDA queries	0%	0%	90%	10%	6	q1.hql ... q7.hql
Stage 2 — EDA charts & storytelling	0%	0%	80%	20%	2	output/q*.jpg, q*.csv
Stage 3 — Feature pipeline	0%	100%	0%	0%	10	ml_data_preparation.py
Stage 3 — Model training & CV	0%	90%	5%	5%	16	ml_models_training.py, models/
Stage 3 — Evaluation	0%	80%	10%	10%	7	ml_evaluation.py, evaluation.csv
Stage 4 — Superset dashboard	10%	10%	70%	10%	6	Public dashboard
Pipeline orchestration (main.sh)	60%	20%	10%	10%	2	main.sh, stage*.sh
Repository documentation	10%	10%	10%	70%	4	README.MD, docs/
Final report (this document)	5%	5%	5%	85%	8	report/report.pdf
Defence presentation	10%	20%	30%	40%	3	presentation/presentation.pdf
QA & reproducibility checks	15%	15%	15%	55%	5	Pipeline re-runs
Total person-hours					92	

Table 13: Contribution table. Per-row percentages sum to 100% across the four team members; the *Avg hrs* column is the team-wide person-hours spent on each task.

References

- [1] Sebastian Garcia, Agustin Parmisano, and Maria Jose Erquiaga. IoT-23: A labeled dataset with malicious and benign IoT network traffic. <http://doi.org/10.5281/zenodo.4743746>, 2020. Version 1.0.0, [Data set], Zenodo.